CMPS 181, Spring 2018, Project 4
Due Thursday, June 7, 11:59pm on Canvas

Project 4: Implementing a Query Engine with the Extension of the Relation Manager

Introduction

In this project, you will first extend the RelationManager (RM) component that you implemented for project 2 so that the RM layer can orchestrate both the RecordBasedFileManager (RBF) and IndexManager (IX) layers when tuple-level operations happen, and the RM layer will also manage new catalog information related to indexes at this level. After the RM layer extension, you will implement a QueryEngine (QE) component. The QE component provides classes and methods for answering SQL queries. For simplicity, you only need to implement several basic relational operators. All operators are iterator-based. To give you a jumpstart, we've implemented two wrapper operators on top of the RM layer that provide file and index scanning. See the Appendix section below for more details.

Use the new Project 4 codebase.zip as your starting point. As usual, we suggest that you fill it in using our Project 3 solution (rather than your own) to do Project 4. See the test cases included in the codebase for examples of how the operators are used.

Part 4.1: RelationManager Extensions

RelationManager

All of the methods that you implemented for Project 2 should now be extended to coordinate data files plus any associated indexes of the data files. For example, if you insert a tuple into a table using RelationManager::insertTuple(), the tuple should be inserted into the table (via the RBF layer) and a corresponding entry should be inserted into each associated index of the table (via the IX layer). Also, if you delete a table using RelationManager::deleteTable(), all associated artifacts (e.g., indexes) should be deleted as well. This applies to both the catalog entries that record information, and to the file artifacts themselves. The RelationManager class, in addition to enforcing the coordination semantics between data files and the indexes in your existing methods, will also include the following newly-added index-related methods. These methods can all be implemented by simply delegating their work to the underlying IX layer that you built in Project 3. (This part of the project is largely a big wrapper.)

```
class RelationManager
{
public:
  ...
  RC createIndex(const string &tableName, const string &attributeName);

  RC destroyIndex(const string &tableName, const string &attributeName);

  // indexScan returns an iterator to allow the caller to go through qualified entries in
index
  RC indexScan(const string &tableName,
                         const string &attributeName,
                         const void *lowKey,
                         const void *highKey,
                         bool lowKeyInclusive,
                         bool highKeyInclusive,
                         RM_IndexScanIterator &rm_IndexScanIterator
        );
  ...
}
```

RC createIndex(const string &tableName, const string &attributeName)

This method creates an index on a given attribute of a given table. It should also put an entry into the catalog describing the index.

[Recall that in Project 2, you created catalog tables for Tables and Columns. Now you'll also need to have a catalog table for Indexes. The Index catalog should contain information about the table that the index is defined on, the attribute of that table that the index is defined on, and the Record-Based File in which the data corresponding to the index is stored.]

This method should return an error if no table the specified name exists, or if there is no attribute on that table with the specified name. All other methods in this assignment should also return suitable errors if the operands they reference don't exist.

RC destroyIndex(const string &tableName, const string &attributeName)

This method destroys an index on a given attribute of a given table. It should also delete the entry for that index from the catalog.

RC indexScan(const string &tableName, const string &attributeName, const void *lowKey, const void *highKey, bool lowKeyInclusive, bool highKeyInclusive, RM_IndexScanIterator &rm_IndexScanIterator)

This method should initialize a condition-based scan over the entries in the open index on the given attribute of the given table. If the scan initiation method is successful, a RM_IndexScanIterator object called rm_IndexScanIterator is returned. (Please see the RM_IndexScanIterator class below.) Once underway, by calling RM_IndexScanIterator::getNextEntry(), the iterator should produce the entries of all records whose indexed attribute key falls into the range specified by the lowKey, highKey, and inclusive flags. If lowKey is NULL, it can be interpreted as -infinity. If highKey is NULL, it can be interpreted as +infinity. The format of the parameter lowKey and highKey is the same as the format of the key in IndexManager::insertEntry().

RM_IndexScanIterator

```
class RM_IndexScanIterator {
 public:
  RM_IndexScanIterator();       // Constructor
  ~RM_IndexScanIterator();      // Destructor

  // "key" follows the same format as in IndexManager::insertEntry()
  RC getNextEntry(RID &rid, void *key);       // Get next matching entry
  RC close();                                 // Terminate index scan
};
                        // Terminate index scan };
```

RC getNextEntry(RID &rid, void *key)

This method should set its output parameters rid and key to be the RID and key, respectively, of the next record in the index scan. This method should return RM_EOF if there are no index entries left satisfying the scan condition. You may assume that RM component clients will not close the corresponding open index while a scan is underway.

RC close()

This method should terminate the index scan.

Part 4.2: Query Engine

Iterator Interface

All of the operators that you will implement in this part inherit from the following **Iterator** interface.

```
class Iterator {
    // All the relational operators and access methods are iterators
    // This class is the super class of all the following operator classes
    public:
        virtual RC getNextTuple(void *data) = 0;
        // For each attribute in vector<Attribute>, name it rel.attr
        virtual void getAttributes(vector<Attribute> &attrs) const = 0;
        virtual ~Iterator() {};
};
```

virtual RC getNextTuple(void *data)

This method should set the output parameter **data** of the next record. The format of the **data** parameter, which refers to the next tuple of the operator's output, is the same as that used in previous projects. Also, null-indicators for the given attributes are always placed at the beginning of **data**. That is, the tuple value is a sequence of binary attribute values in which null-indicators are placed first and then each value is represented as follows: (1) For INT and REAL: use 4 bytes; (2) For VARCHAR: use 4 bytes for the length followed by the characters.

virtual void getAttributes(vector<Attribute> &attrs)

This method returns a vector of attributes in the intermediate relation resulted from this iterator. That is, while the previous method returns the tuples from the operator, this method makes the associated schema information for the returned tuple stream available in the query plan. The names of the attributes in vector<Attribute> should be of the form relation.attribute to clearly specify the relation from which each attribute comes.

Filter Interface

```
class Filter : public Iterator {
    // Filter operator
    public:
        Filter(Iterator *input,                          // Iterator of input R
               const Condition &condition                // Selection condition
        );
        ~Filter();

        RC getNextTuple(void *data);
        // For attribute in vector<Attribute>, name it as rel.attr
        void getAttributes(vector<Attribute> &attrs) const;
};
```

Using this iterator, you can do a selection query such as "SELECT * FROM EMP WHERE sal > 100000".

This filter class is initialized by an input iterator and a selection condition. It filters the tuples from the input iterator by applying the filter predicate **condition** on them. For simplicity, we assume this filter only has a single selection condition. The schema of the returned tuples should be the same as the input tuples from the iterator.

Project Interface

```
class Project : public Iterator {
    // Projection operator
    public:
        Project(Iterator *input,                        // Iterator of input R
                const vector<string> &attrNames);       // vector containing attribute
names
        ~Project();

        RC getNextTuple(void *data);
        // For attribute in vector<Attribute>, name it as rel.attr
        void getAttributes(vector<Attribute> &attrs) const;
};
```

This project class takes an iterator and a vector of attribute names as input. It projects out the values of the attributes in the **attrNames**. The schema of the returned tuples should be the attributes in attrNames, in the order of attributes in the vector.

Index Nested-Loop Join Interface

```
class INLJoin : public Iterator {
    // Index Nested-Loop join operator
    public:
        INLJoin(Iterator *leftIn,                       // Iterator of input R
                IndexScan *rightIn,                     // IndexScan Iterator of input S
                const Condition &condition              // Join condition
        );

        ~INLJoin();

        RC getNextTuple(void *data);
        // For attribute in vector<Attribute>, name it as rel.attr
        void getAttributes(vector<Attribute> &attrs) const;
};
```

The INLJoin iterator takes two iterators as input. The **leftIn** iterator works as the outer relation, and the **rightIn** iterator is the inner relation. The **rightIn** is an object of IndexScan Iterator. Again, we have already implemented the IndexScan class for you, which is a wrapper on RM_IndexScanIterator. The returned schema should be the attributes of tuples from leftIn concatenated with the attributes of tuples from rightIn. You don't need to remove any duplicate attributes.

The only join condition that you need to consider for this assignment is equality (EQ_OP).

Important Note

You must make sure that all operators which create temporary rbfm files clean up after themselves. That is, such files must be deleted when the operator is closed.

An Example

Here is an example showing how to assemble the operators to form query plans. Example: "SELECT Employee.name, Employee.age, Employee.DeptID, Department.Name FROM Employee JOIN Department ON Employee.DeptID = Department.ID WHERE Employee.salary > 50000". We assume for this example that the optimizer has picked the Index Nested-Loop algorithm to execute the join.

```
/****** ****** ****** ****** ******
 *    TABLE SCANS
 ****** ****** ****** ****** ******/

TableScan *emp_ts = new TableScan(rm, "Employee");
TableScan *dept_ts = new TableScan(rm, "Department");

/****** ****** ****** ****** ******
 *    FILTER Employee Table
 ****** ****** ****** ****** ******/

Condition cond_f;
cond_f.lhsAttr = "Employee.Salary";
cond_f.op = GT_OP;
cond_f.bRhsIsAttr = false;
Value value;
value.type = TypeInt;
value.data = malloc(bufsize);
*(int *)value.data = 50000;
cond_f.rhsValue = value;

Filter *filter = new Filter(emp_ts, cond_f);

/****** ****** ****** ****** ******
 *    PROJECT Employee Table
 ****** ****** ****** ****** ******/

vector<string> attrNames;
attrNames.push_back("Employee.name");
attrNames.push_back("Employee.age");
attrNames.push_back("Employee.DeptID");

Project project(filter, attrNames);

/****** ****** ****** ****** ******
 *   INDEX NESTED-LOOP JOIN Employee with Dept
 ****** ****** ****** ****** ******/

Condition cond_j;
cond_j.lhsAttr = "Employee.DeptID";
cond_j.op = EQ_OP;
cond_j.bRhsIsAttr = true;
cond_j.rhsAttr = "Department.ID";

INLJoin *inlJoin = new INLJoin(project, dept_ts, cond_j);


void *data = malloc(bufsize);
while(inlJoin.getNextTuple(data) != QE_EOF)
{
  printAttributes(data);
}
```

Appendix

Below we list the APIs for the three classes used in the operators. For more detailed implementation information, please refer to the **qe.h** header file in the code base. Note that in the TableScan and IndexScan classes, the argument **alias** is used to rename the input relation. In the case of self-joins, at least one of the uses of the relations must be renamed to differentiate the two from each other in terms of attribute naming.

```
struct Condition {
    string lhsAttr;         // left-hand side attribute
    CompOp  op;             // comparison operator
    bool    bRhsIsAttr;     // TRUE if right-hand side is an attribute and not a value;
FALSE, otherwise
    string  rhsAttr;        // right-hand side attribute if bRhsIsAttr = TRUE
    Value   rhsValue;       // right-hand side value if bRhsIsAttr = FALSE
};


class TableScan : public Iterator
{
TableScan(RelationManager &rm, const string &tableName, const char *alias = NULL);
                                    // constructor
void setIterator();
                                        // Start a new iterator
RC getNextTuple(void *data);
                                        // Return the next tuple from the iterator
void getAttributes(vector<Attribute> &attrs) const;
                                        // Return the attributes from this iterator
~TableScan();
                                        // destructor
};


class IndexScan : public Iterator
{
IndexScan(RelationManager &rm, const string &tableName, const string &attrName,
          const char *alias = NULL);
                                        // constructor
void setIterator(void* lowKey, void* highKey, bool lowKeyInclusive, bool highKeyInclusive);
                                        // Start a new iterator given the new compOp and value
RC getNextTuple(void *data);
                                        // Return the next tuple from the iterator
void getAttributes(vector<Attribute> &attrs) const;
                                        // Return the attributes from this iterator
~IndexScan();
                                        // destructor
};
```

Testing

- Please use the tests that we've provided in the codebase to test your code. The qe tests in test.sh use data that we've provided in a data directory that's in the codebase. Our public tests are not an exhaustive test suite. Please feel free to add more cases to this, and test your code thoroughly. We will grade your submission based on a combination of these public tests and our own private tests.

Submission Instructions

The following are requirements on your submission. Points may be deducted if they are not followed.

- Write a report to briefly describe the design and implementation of your query engine module.

- You need to submit the source code under the "rbf" ,"rm" ,"ix", "qe" folders. Make sure you do a "make clean" first, and do NOT include any useless files (such as binary files and data files). Your makefile should make sure the test files provided, *qetest_XX.cc*, **compile and run properly.**

- Please organize your project in the following directory hierarchy: project4-*teamID* / codebase / {rbf, rm, ix, qe, data, makefile.inc, readme.txt, project4-report} where rbf, rm, ix, and qe folders include your source code and makefiles. The data directory should be identical to the data directory that we provided.

- Compress project4-*teamID* into a SINGLE zip file. Each group only submits one file, with the name "project4-*teamID*.zip". (e.g., project4-01.zip)

- Please use the test.sh file that appears under Resources→Lab4 on Piazza to test your code. Put test.sh and your zip file under the same directory. Run test.sh to check whether your project can be properly unzipped and tested, using your own makefile.inc when you are testing the script. If the script doesn't work correctly, it's most likely that your folder organization doesn't meet our requirements.

  Note that this file will be used to grade your project partially since we also have our own private test cases. This is by no means an exhaustive test suite. Please feel free to add more cases to this, and test your code thoroughly. As with previous projects, we will also test your code using a private test cases suite.

- Upload the zip file "project4-*teamID*.zip" to Canvas.

Grading Rubric

Full Credit: 100 points

You must follow the usual submission requirements in order to receive credit for this assignment, but there are no separate points given for this.

- o Code is structured as required
- o Makefile works
- o Code compiles and links correctly

Please make sure that your submission follows these requirements and can be unzipped and built by running the script we provided **on unix.ucsc.edu**. Failure to follow the instructions could make your project unable to build on our machines, which may result in loss of points.

1. Project Report (10%)

   - o Clearly document the design of your project by using the provided template.

2. Implementation details (20%)

   - o Implement all the required functionalities.
   - o Implementation should follow the basic requirements of each function, such as file should be created on disk.

3. Pass the provided test suite (48%)

   - o Each of the 8 public tests may be graded as pass/fail.

4. Pass the private test suite (22%)

   - o There are also some private tests, each of which may be graded as pass/fail.